

Curs 13

Distributed Computing Patterns


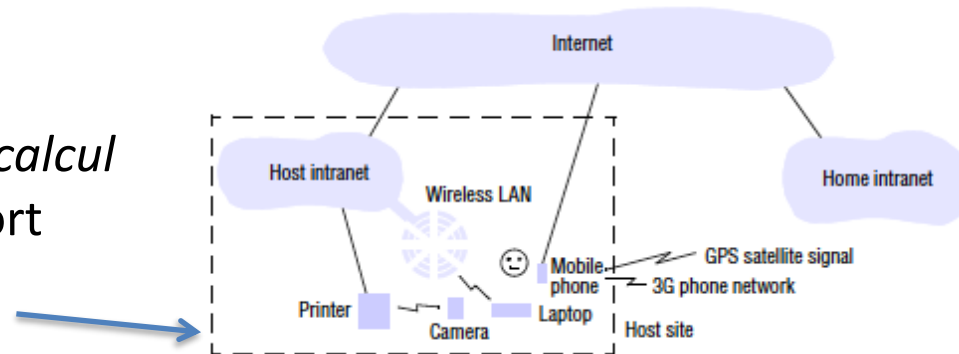
Distributed systems

- Un sistem distribuit poate fi definit ca fiind format din **componente hardware si software localizate intr-o retea de calculatoare care comunica si isi coordoneaza actiunile doar bazat pe transmitere de mesaje**.
- Calcul distribuit (Distributed-Computing) = rezolvarea unor probleme folosind sisteme distribuite

Tendinte

Influente semnificative:

- emergenta tehnologiilor de retea de scara larga
- emergenta *necesaritatilor crescute de calcul* cuplata cu necesitatea de asigura suport pentru *mobilitatea utilizatorilor*
- cresterea cererii de servicii multimedia
- perspectiva asupra sistemelor distribuite ca fiind o utilitate publica



Caracteristici

Concurenta:

- *se lucreaza cu programe care se executa concurent si care partajeaza resurse*

No global clock:

- *nu exista notiunea de timp global*
- aceasta este o consecinta directa a faptului ca singura modalitate de comunicare este transmiterea de mesaje printr-o retea

Independent failures:

Fiecare componenta a sistemului poate 'cadea' in mod independent lasand celelalte in starea de executie ('run').

- toate retelele de calculatoare pot esua('fail') si este responsabilitatea proiectantilor sistemului sa gestioneze efectele in aceste cazuri si sa asigure masuri de reorganizare.
- Aceste tipuri de probleme conduc la izolarea computerelor conectate prin aceste retele aflate in starea de 'fault' – dar nu inseamna si oprirea celorlalte computere;
- Programele pot sa continue sa functioneze dar nu pot detecta daca retea a cazut sau este doar incetinita.
- Similar, 'caderea' unui computer, sau oprirea neasteptata a unui software undeva in sistem (*a crash*), nu este imediat facuta cunoscuta celorlalte componente cu care acesta comunica

Characterization of Faults

- A failure occurs when an error/fault reaches the service interface and alters the service.
- Domain
 - Hardware faults
 - Software faults
- Intent
 - Non-malicious
 - Malicious
- Transient (soft) faults/errors
 - Occurs once and disappears
 - Eg, bit-flip due to high-energy particles
 - Tend to be due to transient physical phenomena
- Intermittent faults/errors
 - Occurs occasionally
 - Eg, a router drops some packets
- Permanent (hard) faults/errors
 - Occurs and doesn't go away
 - Eg, a dead power supply

Examples

- A message-passing application
 - A fix set of N processes
 - Communication by exchanging messages
 - MPI application
 - Cooperate to execute a distributed algorithm
- An asynchronous distributed system
 - Finite set of communication channels connecting any ordered pair of processes
 - Reliable
 - FIFO
 - Ex: TCP, MPI
 - Asynchronous
 - Unknown bound on message transmission delays
 - No order between messages on different channels

Crash failures <-> Fault tolerance


- Crash failures
 - When a process fails, it stops executing and communicating.
 - All data stored locally is lost
- Fault tolerance
 - How to ensure the correct execution of the application in the presence of faults?
 - The execution should terminate
 - It should provide the correct result

Backward error recovery

- Rollback-recovery
 - Restores the application to a previous error-free state when a failure is detected
 - Information about the state of the application saved during failure free execution
 - Assumes the error will be gone when resuming execution
 - Transient (soft) error
 - Use spare resources to replace faulty ones in case of hard error
- BER techniques
 - Checkpointing: saving the system state
 - Logging: saving changes made to the system

Patterns

Client-Server pattern



- O componenta de tip server care furnizeaza servicii catre mai multe componente client.
- O componenta client cere servicii de la componenta server.
- Serverele sunt active permanent 'ascultand' cererile de la clienti.

Starea(State) in sablonul Client-server

Clientii si serverele lucreaza in general in sesiuni ('sessions').

–**stateless server** -- starea unei sesiuni (session state) este gestionata de catre client. Aceasta stare (client) este trimisa impreuna cu fiecare cerere. In aplicatiile web, *session state* poate fi stocata ca si parametrii URL, in campuri ascunse sau folosind cookies (obligatoriu pentru arhitecturile REST folosite pentru aplicatii web).

–**stateful server** -- starea unei sesiuni (session state) este mentinuta la nivel de server si este asociata cu ID clientului


Modalitatea de gestionare a starii unei sesiuni influenteaza tranzactiile, scalabilitatea si gestiunea erorilor (*fault handling*).

- Tranzactiile trebuie sa fie
 - atomice si sa asigure consistenta starii,
 - izolate (sa nu afecteze alte tranzactii)
 - durabile
- *fault handling* => starea mentinuta la nivelul clientului implica faptul ca toata informatia se va pierde in cazul in care clientul esueaza (*fails*).
- Securitatea poate fi afectata daca starea se mentine la nivel de client pentru ca informatia se transmite de fiecare data (la fiecare *request*).
- Scalabilitatea poate fi reduasa daca starea se mentine la nivelul serverului '*in-memory*' – multi client, multe cereri => necesar de memorie crescute.

Peer-to-peer pattern

- Poate fi privit ca si un sablon Client-Server **simetric**:
 - un nod (*peer*) poate functiona ca si un client – care cere servicii de la alte componente sau ca si server care furnizeaza servicii pentru altii.
 - rolul unui nod se poate schimba in mod dinamic
- Atat clientii cat si serverele folosesc in mod uzual multithreading.
- Serviciile pot fi implicite (de exemplu prin intermediul unui stream de conectare) in locul unei cereri (request) trimise prin invocare directa.
- Un *peer* care actioneaza ca si server isi poate informa colegii (peers) care activeaza ca si clienti de aparitia unor evenimente; clientii pot fi informati folosind, de exemplu, o magistrala de evenimente (event-bus).

Example



- peer-to-peer file-sharing like Gnutella or Bittorrent
- the distributed search engine Sciencenet
- multi-user applications like a drawing board

Caracteristici

- **Performanta** creste atunci cand numarul de noduri creste, dar scade atunci cand sunt prea putine

Avantaje

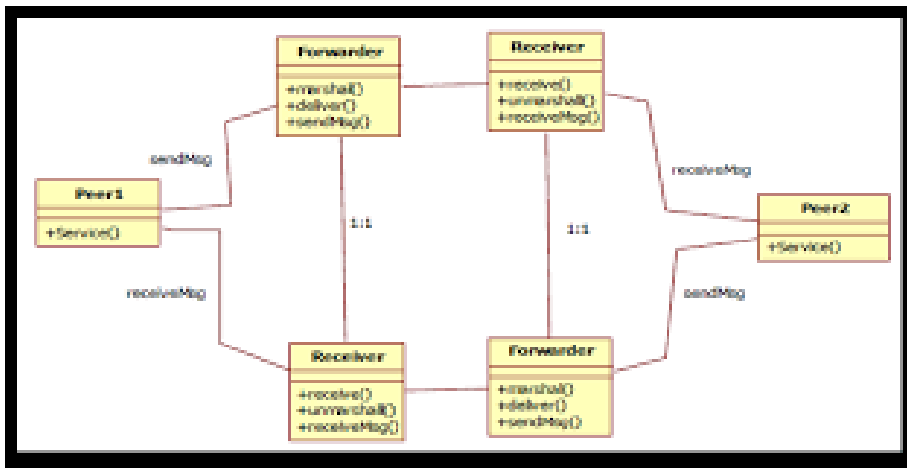
- Nodurile pot folosi capacitatea intregului sistem chiar daca fiecare are o capacitate proprie limitata.
 - este un **cost individual mic cu beneficiu mare** obtinut prin partajare
- Overhead-ul de administrare este scazut pentru ca retelele peer-to-peer se organizeaza intern (self-organizing)
- Asigura **scalabilitate** foarte buna si este rezilienta la esec(failure) componentelor individuale.
- Configurarea sistemului se poate schimba **dinamic**: un *peer* poate intra sau pleca in timp ce sistemul functioneaza.

Dezavantaje

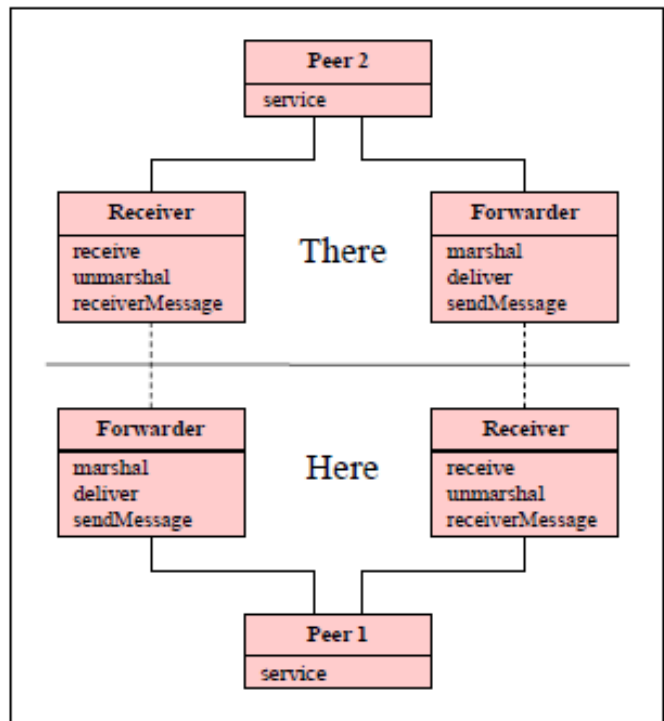
- nu exista garantia calitatii serviciilor (*no guarantee about quality of service*) deoarece nodurile coopereaza voluntar
- nu exista garantia securitatii (*security is difficult to guarantee*) deoarece nodurile coopereaza voluntar

Communication Pattern: Forwarder-Receiver

- Forwarder-Receiver furnizeaza in mod transparent comunicarea inter-proces pentru sistemele sistem cu model de interactiune de tip peer-to-peer.
- Foloseste *forwarders* si *receivers* pentru a decupla nodurile de mecanismul de comunicare de baza (the underlying mechanism).
- *Forwarders* sunt responsabili pentru transmiterea de mesaje folosind mecanisme specifice IPC (Inter Process Communication)
- *Receivers* sunt responsabili pentru receptionarea IPC requests sau mesaje trimise de catre alti peers folosind mecanisme IPC si pentru apelarea metodelor corespunzatoare.



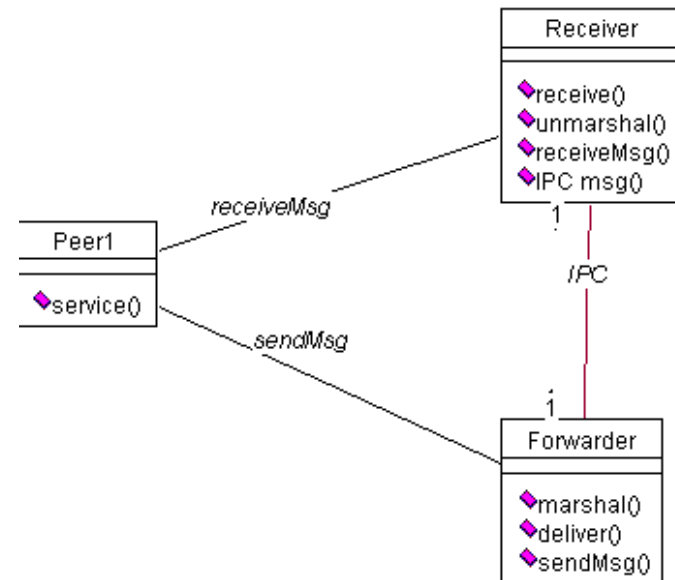
Forwarder-Receiver



- The details of the underlying IPC mechanism for sending or receiving messages are hidden from the peers by encapsulating all system-specific functionality into separate components.
- Examples of such functionality are
 - the mapping of names to physical locations,
 - the establishment of communication channels, or
 - the marshaling and unmarshaling of messages
(marshalling = producing a stream of byte which contain enough information to be able to re-build the object)

F-R - types of messages

- *Command message*- instruct the recipient to perform some activities
- *Information message*- contains data
- *Response message*- allow agents to acknowledge the arrival of a message
 - It includes functionality for sending and marshaling



Forwarder-Receiver

- Peer components are responsible for application tasks.
- Peers may be located on a different process, or even on a different machine.
 - It uses a forwarder to send messages to other peers and a receiver to receive messages from other peers.
 - They continuously monitor network events and resources, and listen for incoming messages from remote agents.
 - Each agent may connect to any other agent to exchange information and requests.
 - To send a message to remote peer, it invokes the method `sendmsg` of its forwarder.
 - It uses `marshal.sendmsg` to convert messages that IPC understands.
 - To receive it invokes `receivemsg` method of its receiver to unmarshal it uses `unmarshal.receivemsg`.
 - Forwarder components send messages across peers.
 - When a forwarder sends a message to a remote peer, it determines the physical location of the recipient by using its name-to-address mapping.
 - Receiver components are responsible for receiving messages.
 - It includes functionality for receiving and unmarshaling messages.

F_R Dynamics

- P1 requests a service from a remote peer P2.
- It sends the request to its forwarder forw1 and specifies the name of the recipient.
- forw1 determines the physical location of the remote peer and marshals the message.
- forw1 delivers the message to the remote receiver recv2.

At some earlier time P2 has requested its receiver recv2 to wait for an incoming requests.

- recv2 receives the message arriving from forw1.
- recv2 unmarshals the message and forwards it to its peer P2.
- Meanwhile P1 calls its receiver recv1 to wait for a response.
- P2 performs the requested service and sends the result and the name of the recipient P1 to the forwarder forw2.
- The forwarder marshals the result and delivers it recv1.
- Recv1 receives the response from P2, unmarshals it and delivers it to P1.

Architectural patterns

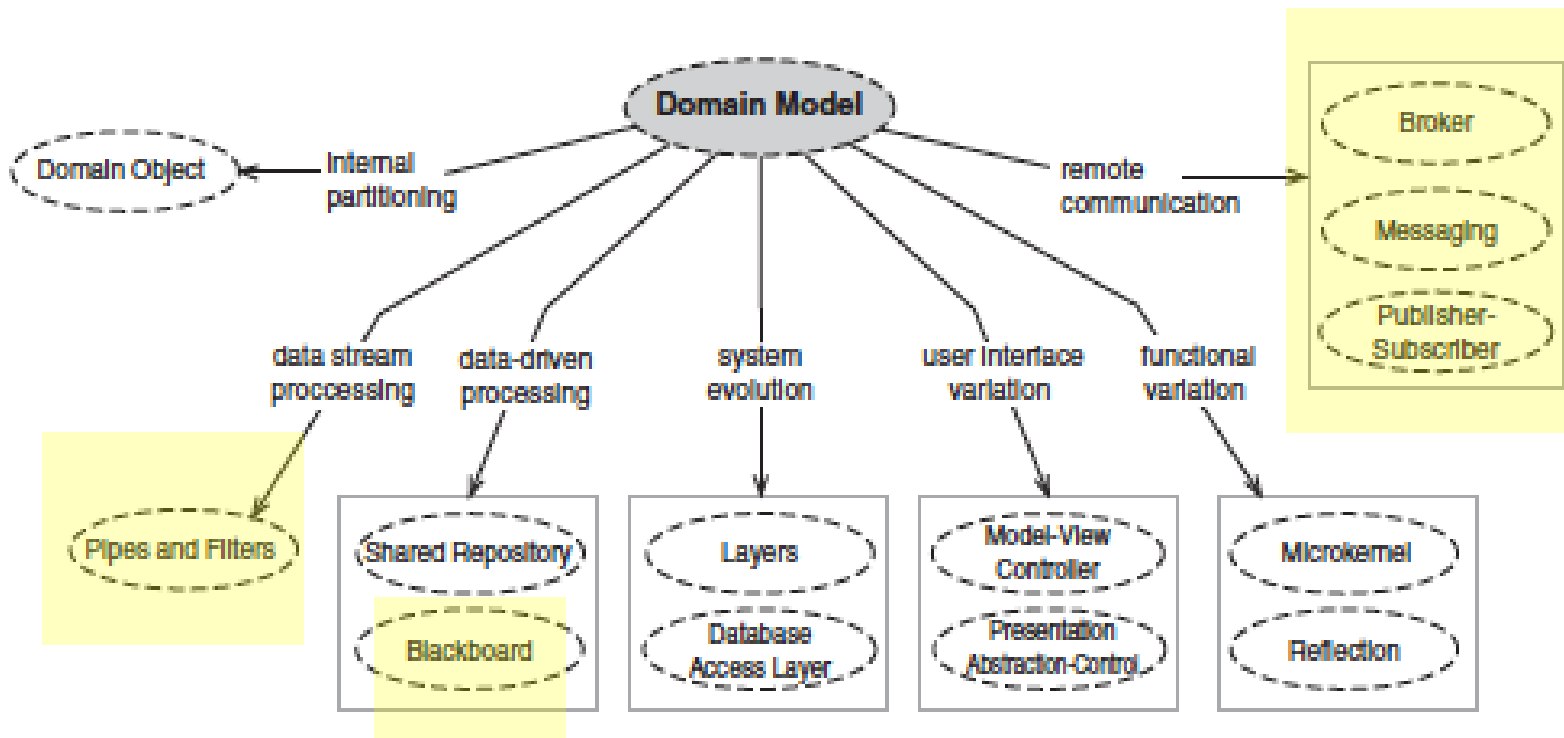
Frank Buschmann. Kevlin Henney. Douglas C. Schmidt. *Pattern-Oriented Software Architecture*, Volume 4:

A Pattern Language for Distributed-Computing. Wiley & Sons, 2007


Un sablon arhitectural este un concept/sablon care
**rezolva si delimiteaza anumite elemente esentiale de coeziune ale unei arhitecturi
software.**

- Domain Model
- Layers
- Model-View-Controller
- Presentation-Abstraction-Control
- Microkernel
- Reflection
- Pipes and Filters
- Shared Repository
- Blackboard
- Domain Object

Conexiunea cu Domain Layer




Pipe-Filter Pattern



- furnizeaza pentru un sistem o structura care produce un ***stream de date***
- fiecare pas de procesare este incapsulat intr-o componenta de tip *filter*
- Datele sunt transferate prin pipes
 - *pipe* – leaga 2 filtre
 - *pipes* pot fi utilizate pentru sincronizare sau pentru buffering

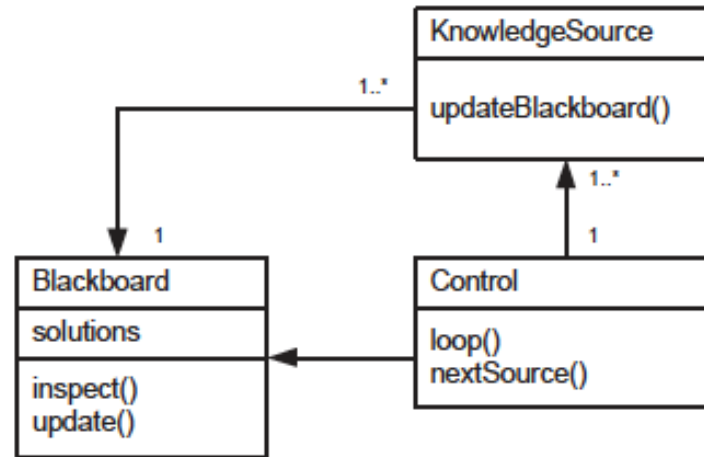
Exemplu



- se poate utiliza pentru a divide taskuri de procesare mari intr-o secventa de componente independente mai mici de procesare (Filters) care sunt conectate prin canale (Pipes).
- fiecare filtru expune o interfata simpla – primeste mesaje de la *inbound pipe*, proceseaza mesajul si publica rezultatul la outbound pipe.
- pentru ca toate componentele utilizeaza aceeasi interfata externa ele pot fi compuse in diferite solutii prin conectarea componentelor la pipe-uri diferite
- se pot adauga filtre sau rearanja in secvente noi fara a se schimba filtrele (in interior).

Blackboard Pattern

sugereaza o modalitate de a rezolva o problema **complexa** (e.g. recunoasterea de imagini sau de limbaj) prin impartirea in subsisteme mai mici specializate care pot rezolva problema **impreuna**.



- mai multe subsisteme specializate asambleaza cunostinte pentru a construi partial sau aproximativ solutia
- ***Toate componentele au acces la ansamblul de date partajate = the blackboard.***
- Componentele pot produce date noi care sunt adaugate pe blackboard.
- Componentele cauta anumite date pe blackboard, folosind e.g. pattern matching.

Componente

- Class
 - **Blackboard**
 - Responsibility
 - Manages central data
 - Collaborators
 - no
- Class
 - **Knowledge Source**
 - Responsibility
 - Evaluates its own applicability
 - Computes a result
 - Updates Blackboard
 - Collaborators
 - Blackboard
- Class
 - **Control**
 - Responsibility
 - Monitors Blackboard
 - Schedules knowledge source activations
 - Collaborators
 - Blackboard
 - Knowledge Source

Descriere generala

- Blackboard permite mai multor procese (agenti) sa comunice prin citirea si scrierea de cereri si informatii catre un depozit de date globale.
- **Fiecare agent participant are o expertiza in propriul domeniu**, si are un anumit tip de cunoastere referitoare la rezolvarea problemei (*knowledge source*) care se poate aplica la o parte a problemei, i.e., problema nu se poate rezolva de catre un singur agent.
- **Agentii comunica strict prin intermediul unei *common blackboard*** care este vizibil tuturor agentilor.
- Atunci cand o problema partiala trebuie rezolvata, potentialii agenti care ar putea sa o rezolve sunt listati.
- O **unitate de control** este responsabila pentru selectarea agentilor si atribuirea de sarcini acestora.

Exemplu: speech recognition

Ciclul principal de rezolvare

- **Control** calls **nextSource()** to select the next knowledge source
- **nextSource()** looks at the blackboard and determines which knowledge sources to call
- For example, **nextSource()** determine that **Segmentation**, **Syllable Creation** and **Word Creation** are candidate
- **nextsource()** invokes the **condition-part** of each candidate knowledge source
- The **condition-parts** of candidate knowledge source inspect the blackboard to determine if and how they can contribute to the current state of the solution
- The **Control** chooses a knowledge source to invoke and a **set of hypotheses** to be worked on (according to the result of the condition parts and/or control data)
- Apply the **action-part** of the knowledge source to the **hypothesis**
- ***New contents are updated in the blackboard***

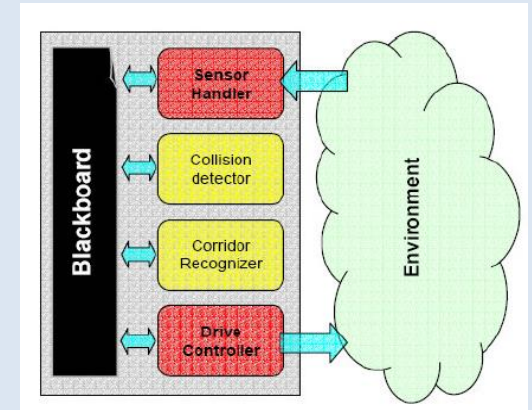
Robot example

- An Experimental robot is equipped with four agents:

- Sensor Handler Agent,
- Collision Detector Agent,
- Corridor Recognizer Agent and
- Drive Controller Agent

(Includes the control software)


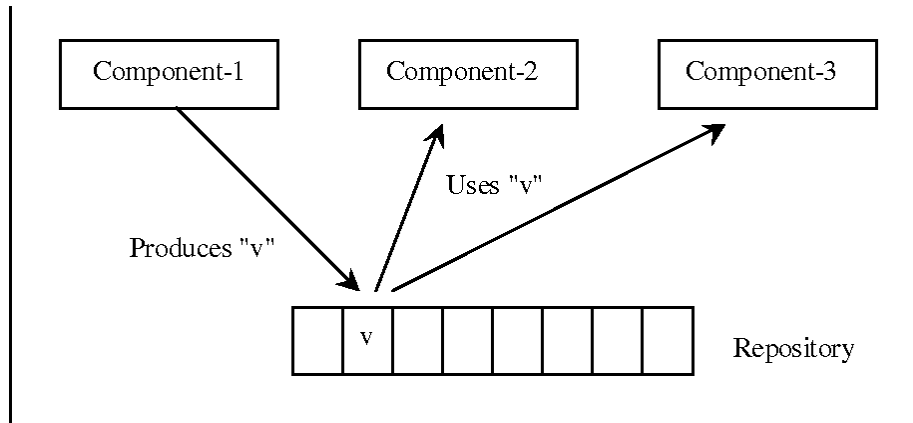
- Agents and blackboard form the control system. Agent cooperation is reached by means of the blackboard. Blackboard is used as a central repository for all shared information.
- Only two agents have access to the environment: Sensor Handler Agent and Drive Controller Agent.
- There is no global controller for Basically, these agents, so each of them independently tries to make a contribution to the system during the course of navigation.
- Basically, each of the four agents executes its tasks independently using information on the blackboard and posts any result back to the blackboard.



Avantaje & Dezavantaje

- Avantaje
 - potrivit pentru diverse surse de date/procesare
 - potrivit pentru medii distribuite
 - **potrivit pentru planificarea taskurilor si a deciziilor**
 - **potrivit pentru abordari de rezolvare in echipa** – se posteaza subcomponente si rezultate partiale
 - util pentru probleme pentru care nu sunt cunoscute strategii de rezolvare deterministe => **opportunistic problem solving**.
- Dezavantaje
 - Scump
 - Dificil de a determina partitionarea pentru knowledge
 - Control unit poate fi foarte complexa

More General: Shared repository



....sabioane legate de infrastructura de comunicare

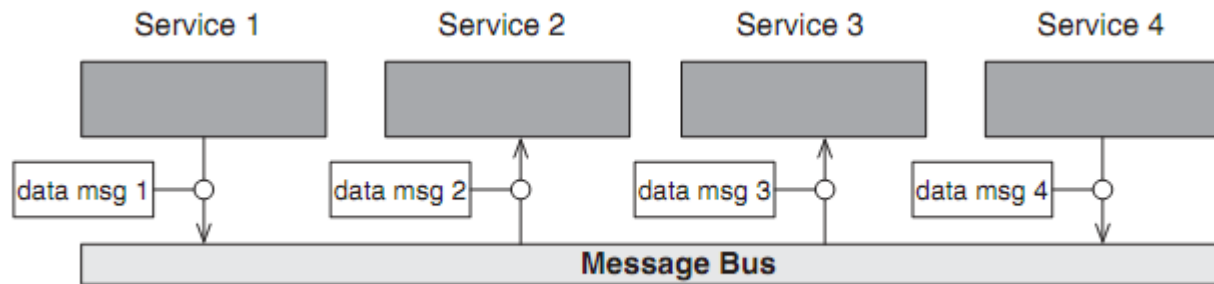
- **Messaging**
 - Distribution Infrastructure
- **Publisher-Subscriber**
 - Distribution Infrastructure
- **Broker**
 - Distribution Infrastructure
- **Client Proxy**
 - Distribution Infrastructure
- **Reactor**
 - Event Demultiplexing and Dispatching
- **Proactor**
 - Event Demultiplexing and Dispatching

Messaging Pattern


- Unele sisteme distribuite sunt compuse din servicii care sunt dezvoltate independent.
- Pentru a forma un sistem coerent aceste servicii trebuie sa interactioneze intr-un mod fiabil fara a implica totusi o dependenta stransa intre ele.
- **Solutie:** Conectarea printr-un canal de mesaje (*a message bus*) care sa permita transferul de date ***asincron***.
 - mesajele se codifica astfel incat comunicarea sa se poata face fara sa fie nevoie de sa fie cunoscute toate informatiile legate de tipurile de date.

Messaging

- Message-based communication suporta *loose coupling*



- Mesajele* contin doar datele care pot fi interschimbate intre setul de clienti si servicii – nu si cine este interesat de acestea.
- => conectarea clientilor si serviciilor printr-un canal care permite schimbul de mesaje “Message Channel”.

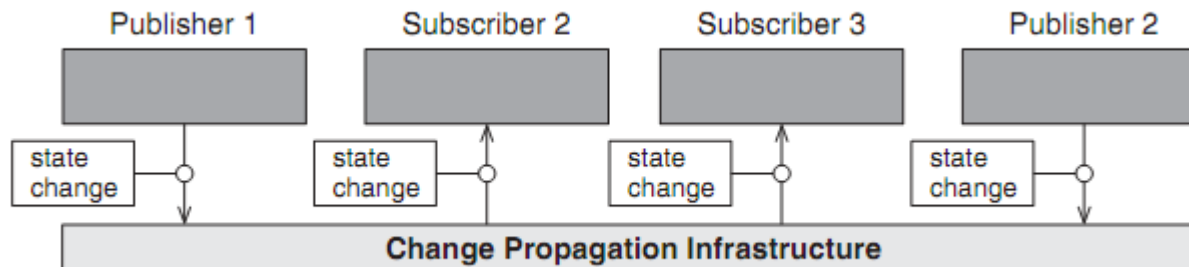


Publisher-Subscriber Pattern

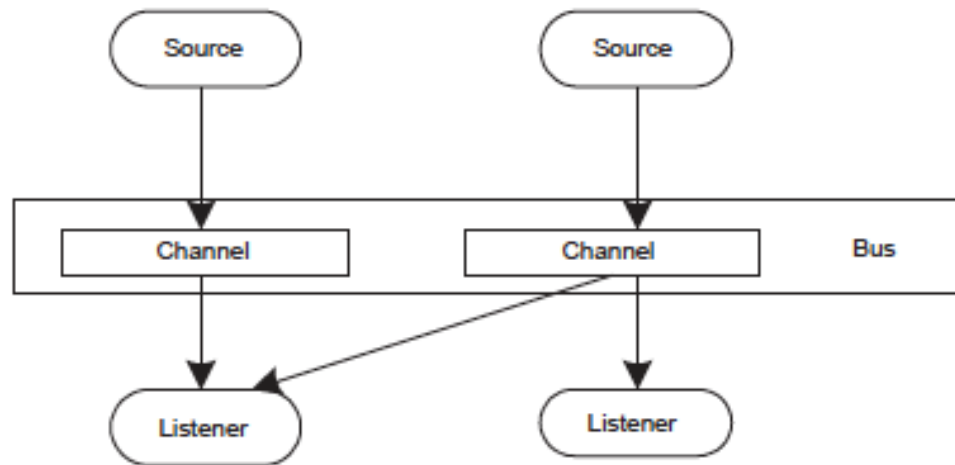
- Componentele din anumite aplicatii distribuite sunt cuplate slab si opereaza in general independent.
- pentru a propaga informatii in asemenea aplicatii se poate folosi un mechanism de notificare prin care sa se faca informari referitoare la modificari de stare sau referitoare la aparitia unor evenimente.
- **Solutie:** Definirea unei infrastructuri de propagare a informatiei care permite editorilor sa disemineze evenimente care contin informatie care ar putea sa intereseze alti actori.
 - se notifica abonatii inregistrati pentru a primi anumite tipuri de notificari

Publisher-Subscriber

- Editorii inregistreaza ce tipuri de evenimente publica.
- Abonatii inregistreaza de ce tipuri de evenimente sunt interesati.
- Infrastructura foloseste informatiile inregistrate pentru a transmite prin retea evenimentele de la editori la abonatii inregistrati a fi interesati.



Event-Bus Pattern




- Sursele de evenimente (*Event sources*) publica mesaje pe canale particulare pe un (*event bus*)
- Event listeners se aboneaza la anumite canale.
- Listenerii sunt notificati referitor la canalele care sunt publicate pe un canal la care s-au abonat.
- Generarea si notificarea de mesaje este **asincrona** :
 - un event source genereaza un mesaj – nu asteapta ca acesta sa fie primit de catre listeneri

Broker Pattern

- Sistemele distribuite pot avea dificultati (provocari) care nu apar in sistemele cu un proces
 - Important – codul aplicatiile nu trebuie sa trateze aceste probleme in mod direct => intermediari (brokers)
- aplicatiile trebuie simplificate prin modularizare *modular programming model* care poate ascunde detaliile de retea si locatie
- **Solutie:** Folosirea unei federatii de brokeri
 - *brokers to separate and encapsulate the details of the communication infrastructure* in a distributed system *from its application functionality*.
- Definirea unui model de tip *component-based programming model prin care clientii sa poate sa invoce metode sau servicii remote ca si cum ar fi locale*

Broker



- .

Client-Proxy Pattern

- Atunci cand se construiește o infrastructura de tip broker *client-side* pentru o componenta *remote* trebuie sa se asigure o abstractizare care permite clientilor sa acceseze acele componente folosind ***remote method invocation***.
- un “Client Proxy / Remote Proxy” reprezinta o componenta remote- in spatiul de adrese al clientului.
- Proxy ofera o interfata identica care mapeaza invocarile de metode specifice catre functionalitatea orientate catre transmiterea de mesaje ale brokerului.
- ***Proxies allow clients to access remote component functionality as if they were collocated.***

Event Demultiplexing & Dispatching

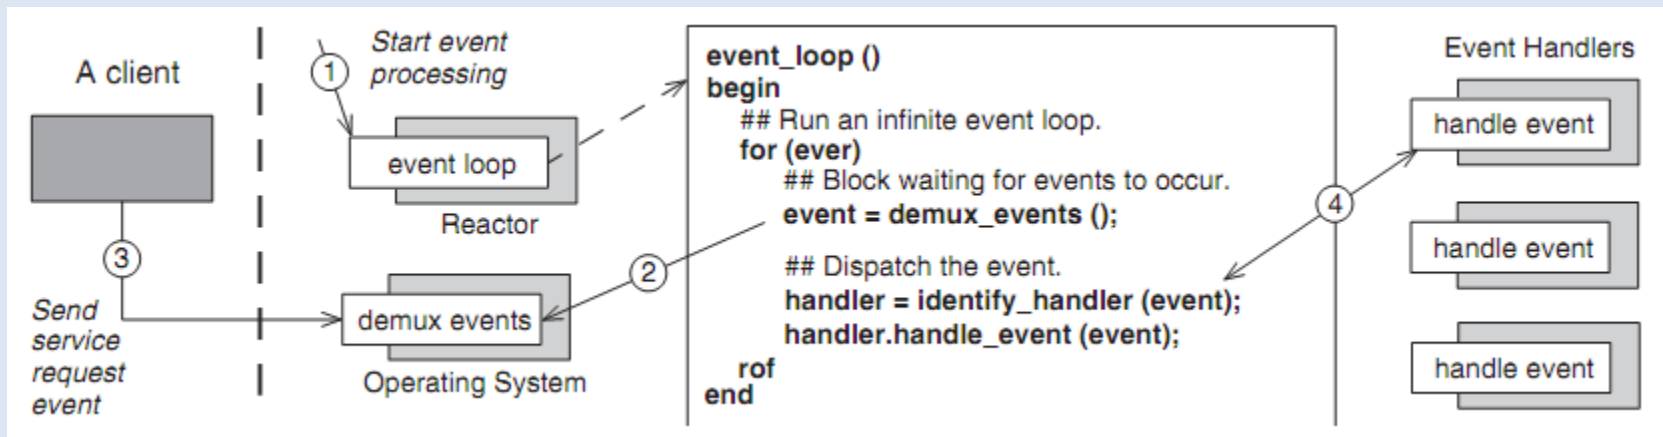
- At its heart, distributed computing is about handling and responding to events received from the network.
- There are patterns that describe different approaches for initiating, receiving, demultiplexing, dispatching, and processing events in distributed and networked systems.
- ...two of these patterns:
 - Reactor ... synchronous
 - Proactor ... asynchronous

Reactor Pattern

- Event-driven software often
 - receives service request events from multiple event sources, which it demultiplexes and dispatches to event handlers that perform further service processing.
- Events can also arrive simultaneously at the event-driven application.
 - To simplify software development, events could be processed sequentially | synchronously.

Reactor

- **Solution:** *Provide an event handling infrastructure that waits on multiple event sources simultaneously for service request events to occur, but only demultiplexes and dispatches **one event at a time** to a corresponding event handler that performs the service.*



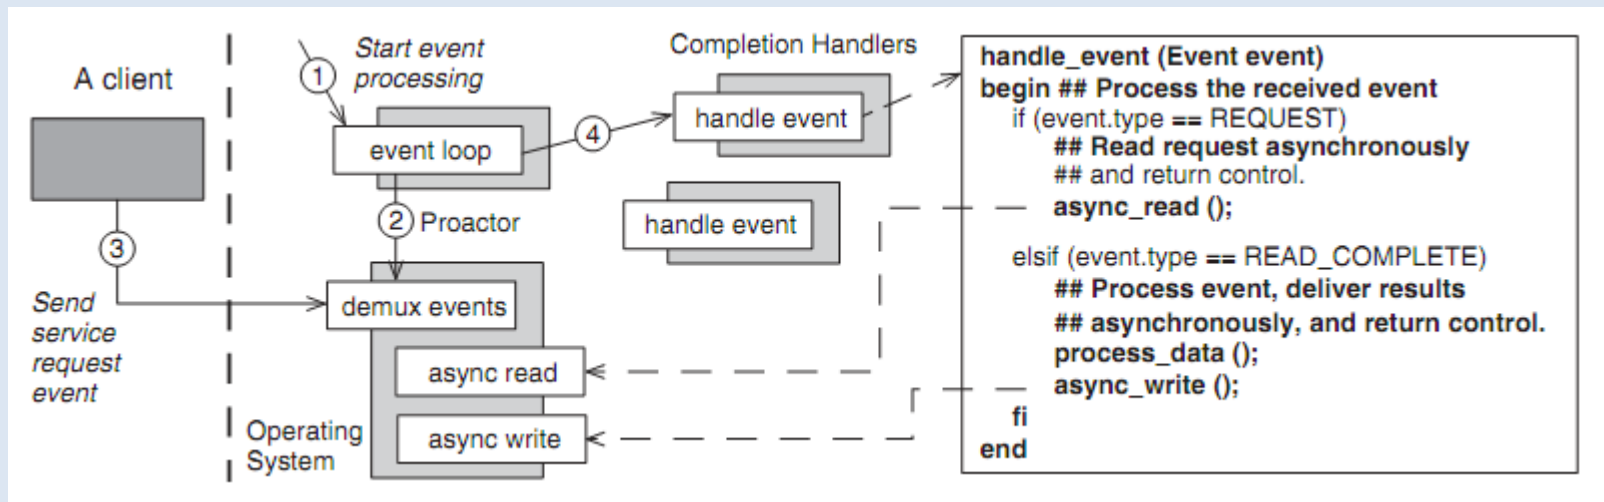
- It defines an **event loop** that uses an **operating system event demultiplexer** to wait **synchronously** for service request events.
- By **delegating the demultiplexing of events to the operating system**, the reactor can wait for multiple event sources simultaneously without a need to multi-thread the application code.

Proactor Pattern

- To achieve the required performance and throughput, event-driven applications must often be able *to process multiple events simultaneously*.
- However, **resolving this problem via multi-threading, may be undesirable, due to the overhead of synchronization, context switching and data movement.**
- **Solution:**
 - *Split an application's functionality into*
 - *asynchronous operations* that perform activities on event sources and
 - *completion handlers* that *use the results of asynchronous operations to implement application service logic*.
 - Let the operating system execute the asynchronous operations, but
 - execute *the completion handlers in the application's thread of control*.


Proactor

- A proactor component coordinates the collaboration between completion handlers and the operating system.
 - It defines an **event loop** that uses an **operating system event demultiplexer** to wait synchronously for events that indicate the completion of asynchronous operations to occur.



- Initially **all completion handlers 'proactively' call an asynchronous operation to wait for service request events to arrive**, and then run the event loop on the proactor.

Proactor




- When such an event arrives, the proactor dispatches the result of the completed asynchronous operation to the corresponding completion handler.
- This handler then continues its execution, which may invoke another asynchronous operation.

Reactor vs. Proactor

- Although both patterns resolve essentially the same problem in a similar context, and also use similar patterns to implement their solutions, the concrete event-handling infrastructures they suggest are distinct, due to the orthogonal forces to which each pattern is exposed.
- REACTOR focuses on simplifying the programming of event-driven software.
 - It implements a *passive* event demultiplexing and dispatching model in which services wait until request events arrive and then react by processing the events synchronously without interruption.
 - While this model scales well for services in which the duration of the response to a request is short, it can introduce performance penalties for long-duration services, since executing these services synchronously can unduly delay the servicing of other requests.

Reactor vs. Proactor

- PROACTOR is designed to maximize event-driven software performance.
 - It implements a more *active* event demultiplexing and dispatching *model* in which services divide their processing into multiple self-contained parts and
 - *proactively initiate asynchronous execution* of these parts.
 - This design allows multiple services to execute concurrently, which can increase quality of service and throughput.
- REACTOR and PROACTOR are not really equally weighted alternatives, but rather are *complementary patterns* that trade-off programming simplicity and performance.
 - Relatively simple event-driven software can benefit from a REACTOR-based design, whereas
 - PROACTOR offers a more efficient and scalable event demultiplexing and dispatching model.



Middleware

- In computer science, a middleware is a software layer that resides between the application layer and the operating system.
- Its primary role is to bridge the gap between application programs and the lower-level hardware and software infrastructure, to coordinate how parts of applications are connected and how they inter-operate.
- Middleware also enables and simplifies the integration of components developed by different technology suppliers.

Middleware

Examples of a middleware for distributed object-oriented enterprise systems: CORBA, SOAP.

- Despite their detailed differences, middleware technologies typically follow one or more of three different communication styles:
 - Messaging
 - Publish/Subscribe
 - Remote Method Invocation

Referinte

- Frank Buschmann. Kevlin Henney. Douglas C. Schmidt. **Pattern-Oriented Software Architecture**, Volume 4: A Pattern Language for Distributed-Computing. Wiley & Sons, 2007
- George Coulouris. Jean Dollimore. Tim Kindberg. Gordon Blair. **DISTRIBUTED SYSTEMS: Concepts and Design**. Fifth Edition. Addison-Wesley.